

De Data:

The One Where We Do De Data

Lecture 2: Text Analytics for Big Data
Mark Keane, Insight/CSI, UCD

Selling
Things

stock-
markets

social
media

science

news

polls

sentiment-id

sentiment-use

time-series

summaries

VSMs

Classifiers

Clustering

cosine

jaccard

dice

levenschtein

TF-IDF

LLR

PMI

Entropy

simple frequencies

*** pre-processed text items of some sort... ***

The Basic View

- ◆ Well, the course is called Text Analytics
- ◆ One expects the *data* to be text, words, clauses, sentences, paragraphs, and other things written in *natural language*
- ◆ Here, we consider this data and *pre-processing* “raw” text for future use

What's the Problem?

“It was in the winter of '69 that we encountered the first sign that the U.S.A. was entering into a conflict, streked (sic) by the GOP, against some unspecified Evil. Mother feared the worst for future fish catches. The Coalition of Baddies was afraid of war and fishing immediately ceased in the South China Sea.”

What's the Problem?

“It was in the winter of '69 that we encountered the first sign that the U.S.A. was entering into a conflict, streked (sic) by the GOP, against some unspecified Evil. Mother feared the worst for future fish catches. The Coalition of Baddies was afraid of war and fishing immediately ceased in the South China Sea.”

What are these ?

What's the Problem?

“It was in the winter of '69 that we encountered the first sign that the U.S.A. was entering into a conflict, **streked** (sic) by the GOP, against some unspecified Evil. Mother feared the worst for future fish catches. The Coalition of Baddies was afraid of war and fishing immediately ceased in the South China Sea.”

This is really “stoked”

What's the Problem?

“It was in the winter of ’69 that we encountered the first sign that the U.S.A. was entering into a conflict, streked (sic) by the GOP, against some unspecified Evil. Mother **feared** the worst for future fish catches. The Coalition of Baddies was **afraid** of war and fishing immediately ceased in the South China Sea.”

“**feared**” & “**afraid**” are sort of the same !

What's the Problem?

“It was in the winter of ’69 that we encountered the first sign that the U.S.A. was entering into a conflict, streked (sic) by the GOP, against some unspecified Evil. Mother feared the worst for future **fish** catches. The Coalition of Baddies was afraid of war and **fishing** immedately ceased in the South China Sea.”

“**fish**” & “**fishing**” look the same but are different !

Not all words are equal...

“winter sign U.S.A. conflict GOP Evil Mother
feared **fish** catches Coalition Baddies afraid
war **fishing** ceased South China Sea”

find the content-bearing words

Will end with...

“winter encounter sign usa enter conflict stok
gop region east asia mother fear future fish
catch coalition bad fear afraid fish cease
south china sea.”

Pre-processing does this; all words are
modified, stemmed, dropped and the output
is the raw data (“words”) used in the analysis

Overview

- ◆ Basics of Natural Language Processing (NLP)
- ◆ Pre-processing to modify text:
 - ◆ tokenisation, stemming, POS tagging, lemmatisation, fixing spellings
- ◆ Pre-processing to exclude (some) text
 - ◆ removing stop words
- ◆ When pre-processing helps or not ?

Important Point

- ◆ Pre-processing is not just about taking things out; stripping off stems, removing stops etc...
- ◆ It may also be about putting things in; like POS tags, syntax, entity tags, lexical chains

NLP 101

I'll -ology you

- ◆ NLP is about meaning, parsing stops short:
 - ◆ *Phonology & Morphology*: sounds & words
 - ◆ *Syntax*: traditionally syntactic constituents
 - ◆ *Meaning*: Semantics & Pragmatics

I'll -ology you.

- ◆ Natural Language Processing concerns itself with recovering meaning from spoken sounds or written strings, usually using different levels of analysis:
 - ◆ Phonology (speech alone)
 - ◆ Morphology (speech & writing)
 - ◆ Syntax (speech & writing)
 - ◆ Semantics (speech & writing)
 - ◆ Pragmatics (speech & writing)

I'll -ology you..

- ◆ *Phonology*: analysis of sounds into phonemes, sets of which form words
- ◆ *Morphology*: analysis of sounds/signs into morphemes (smallest grammatical unit of a language) not= word, morphemes may or may not stand alone ("dog" & "s" versus "dog" & "dogs")
- ◆ *Syntax*: analysis of rules of combination of grammatical units of a language to form sentences

I'll -ology you...

- ◆ *Semantics*: analysis of how the signs of a language come to convey its meaning; the study of meaning at the levels of words, phrases, sentences, and larger units of discourse (termed texts or narratives)
- ◆ *Pragmatics*: analysis of how context contributes to meaning (*yeah...yeah*)

Lemma (morphology)

From Wikipedia, the free encyclopedia

In morphology and lexicography, a **lemma** (plural *lemmas* or *lemmata*) is the **canonical form**, **dictionary form**, or **citation form** of a set of words (**headword**)^[citation needed]. In English, for example, *run*, *runs*, *ran* and *running* are forms of the same **lexeme**, with *run* as the lemma. **Lexeme**, in this context, refers to the set of all the forms that have the same meaning, and **lemma** refers to the particular form that is chosen by convention to represent the lexeme. In lexicography, this unit is usually also the **citation form** or **headword** by which it is indexed. Lemmas have special significance in highly **inflected languages** such as **Turkish** and **Czech**. The process of determining the **lemma** for a given word is called **lemmatisation**. The lemma can be viewed as the chief of the **principal parts**, although lemmatisation is at least partly arbitrary.

- ◆ Is on aspects of morphology and syntax to find...
- ◆ ...parts of speech for words (fish is a noun, fishing is a verb, fishy is an adjective)
- ◆ ...roots and / or lemmas of words (the root fish~ from fishing, fishes, fished; the root be~ from am, are, is)

Focus on Parsing...

Parse - Merriam-Webster Online

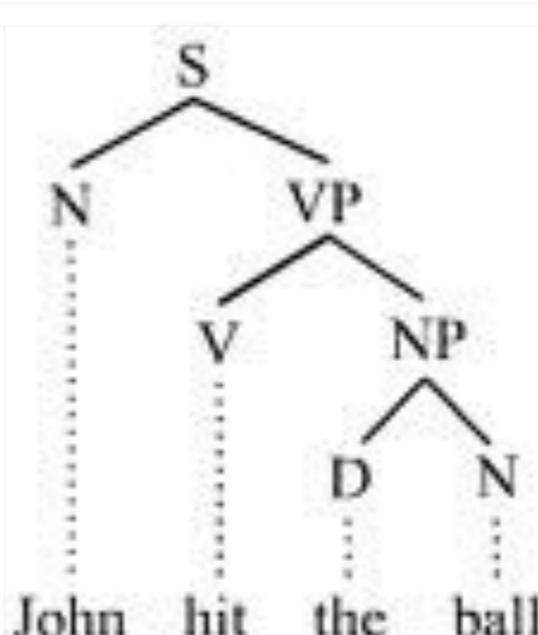
www.merriam-webster.com/dictionary/parse ▾

grammar : to divide (a sentence) into grammatical parts and identify the parts and their relations to each other. : to study (something) by looking at its parts ...

Parsing

Programming Language

Parsing or syntactic analysis is the process of analysing a string of symbols, either in natural language or in computer languages, according to the rules of a formal grammar. The term parsing comes from Latin *pars*, meaning part. [Wikipedia](#)



Constituency-based parse tree

Text Pre-Processing

Pre-Processing: De Guts

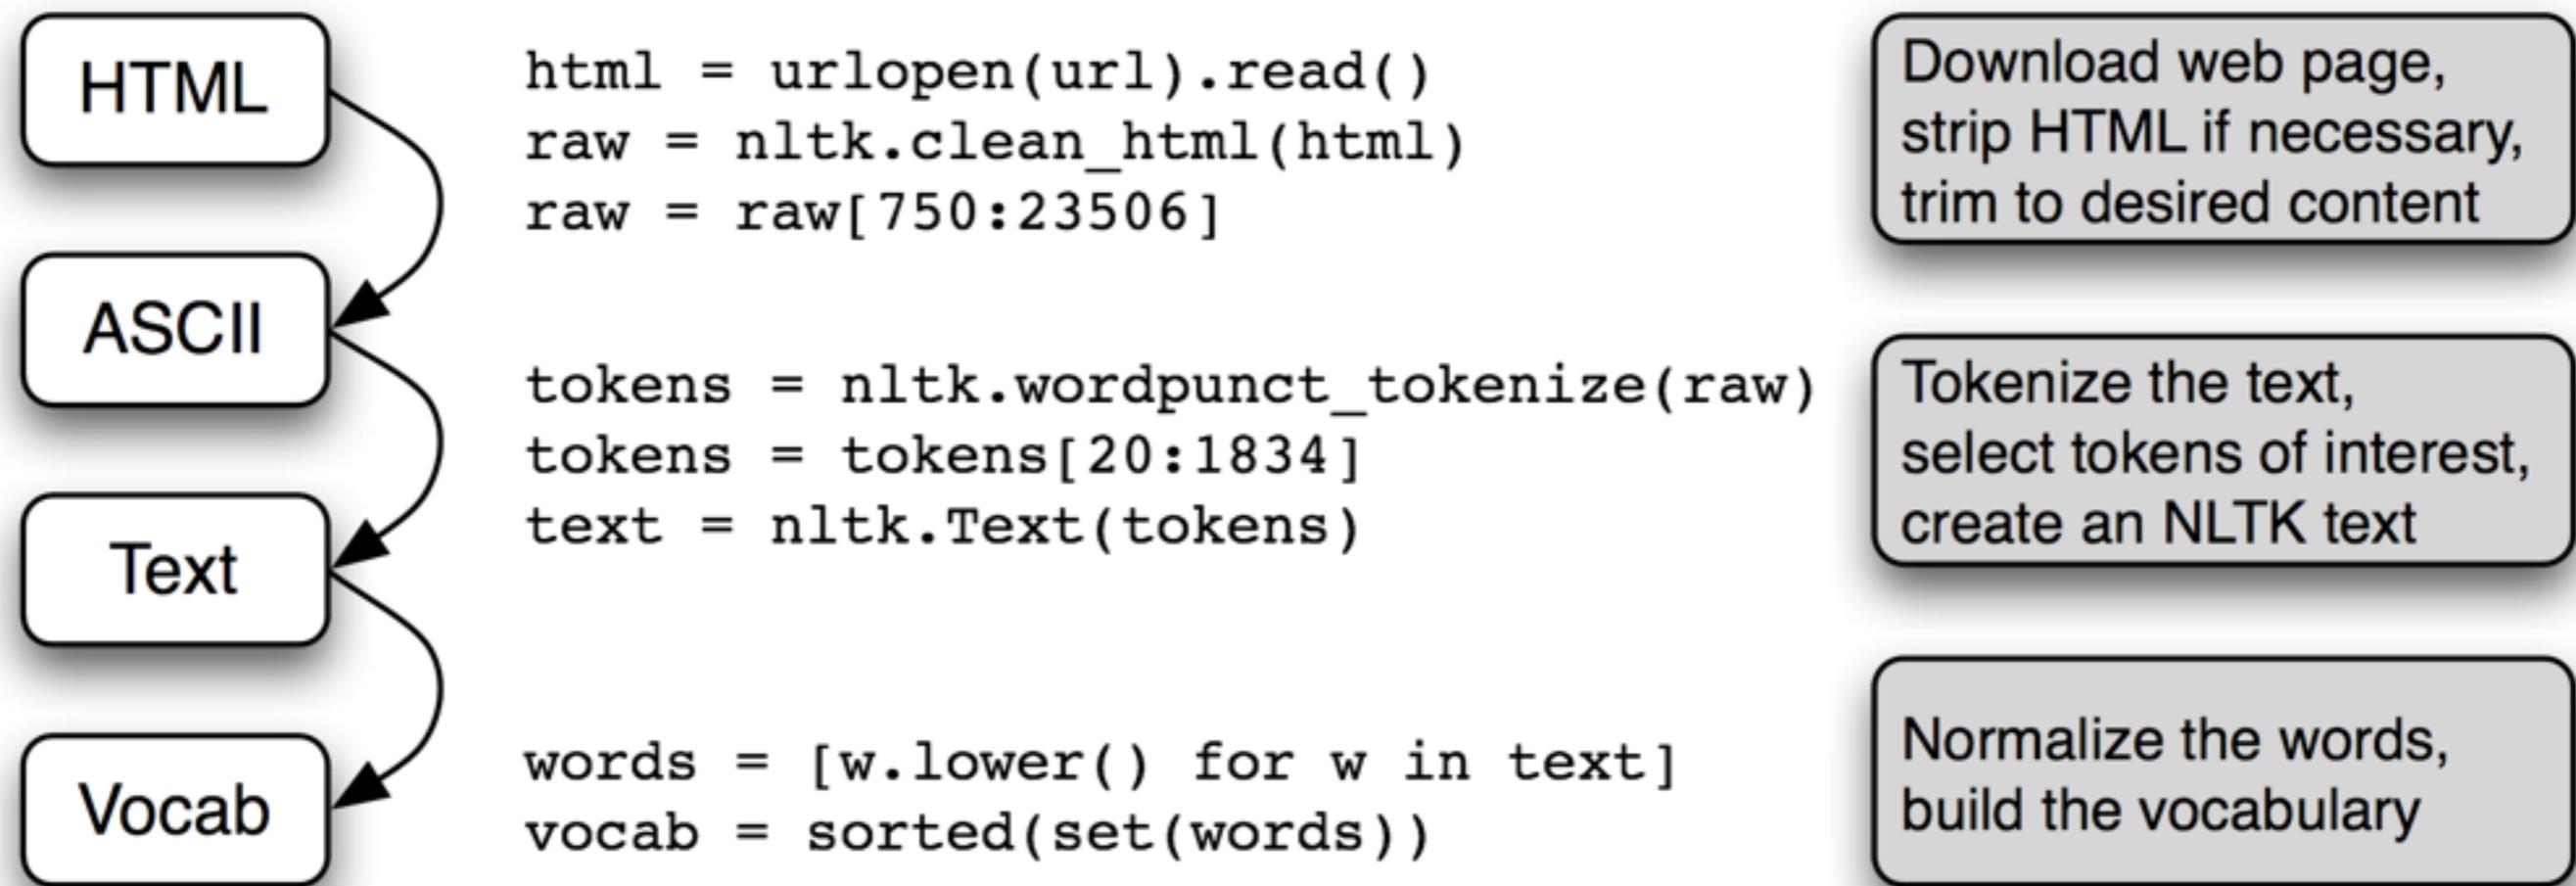
Our Focus...

- ◆ Text pre-processing is the poor-farmer cousin of full NLP; its not really about meaning
- ◆ Its about cleaning up text-data for future use
- ◆ Uses ideas from NLP (eg syntactic analysis, parsing) ... but is not often full NLP
- ◆ Ultimately, it seldom recovers meaning

Standard Tasks

- ◆ *Tokenisation & Normalisation*: finding boundaries between word-like entities in character string
- ◆ *Fixing Misspellings*: where possible
- ◆ *Stemming, lemmatisation, POS-tagging*: finding slightly deeper identities between words (fished, fishing)
- ◆ *Removing Stop Words*: maximising the content-full words in the document/corpus
- ◆ *Entity Extraction*: identifying conceptual entities behind words

Typical Python Pipeline



But...First Remember

- ◆ *Document encodings*: byte sequence turned into a character sequence, ok if ASCII but are different encoding schemes (Unicode UTF-8); need to identify (meta-data) and decode
- ◆ *Document formats*: DOC, PDF, XML, HTML may all need own converters to ensure characters are identified
- ◆ *Document text*: text-part of particular doc needs to be extracted (e.g., from XML, often not clean for sites)

Text Pre-Processing

Tokenisation & Normalisation

What is the Problem ?

- ◆ We have some sequence of characters, we need to break into word-like tokens / entities
- ◆ No problemo, marko, you pick those things with space-chars around them or fullstops before them...
- ◆ Right?

Why Tokenise?

- ◆ Well, what about:

I've: 'I' 've' / 'I've' / Ive

U.S.A.: 'U' 'S' 'A' / 'USA' / 'U.S.A.'

Jonny O'Neill: 'Jonny' 'O' 'Neill' / 'Jonny' 'O'Neill'

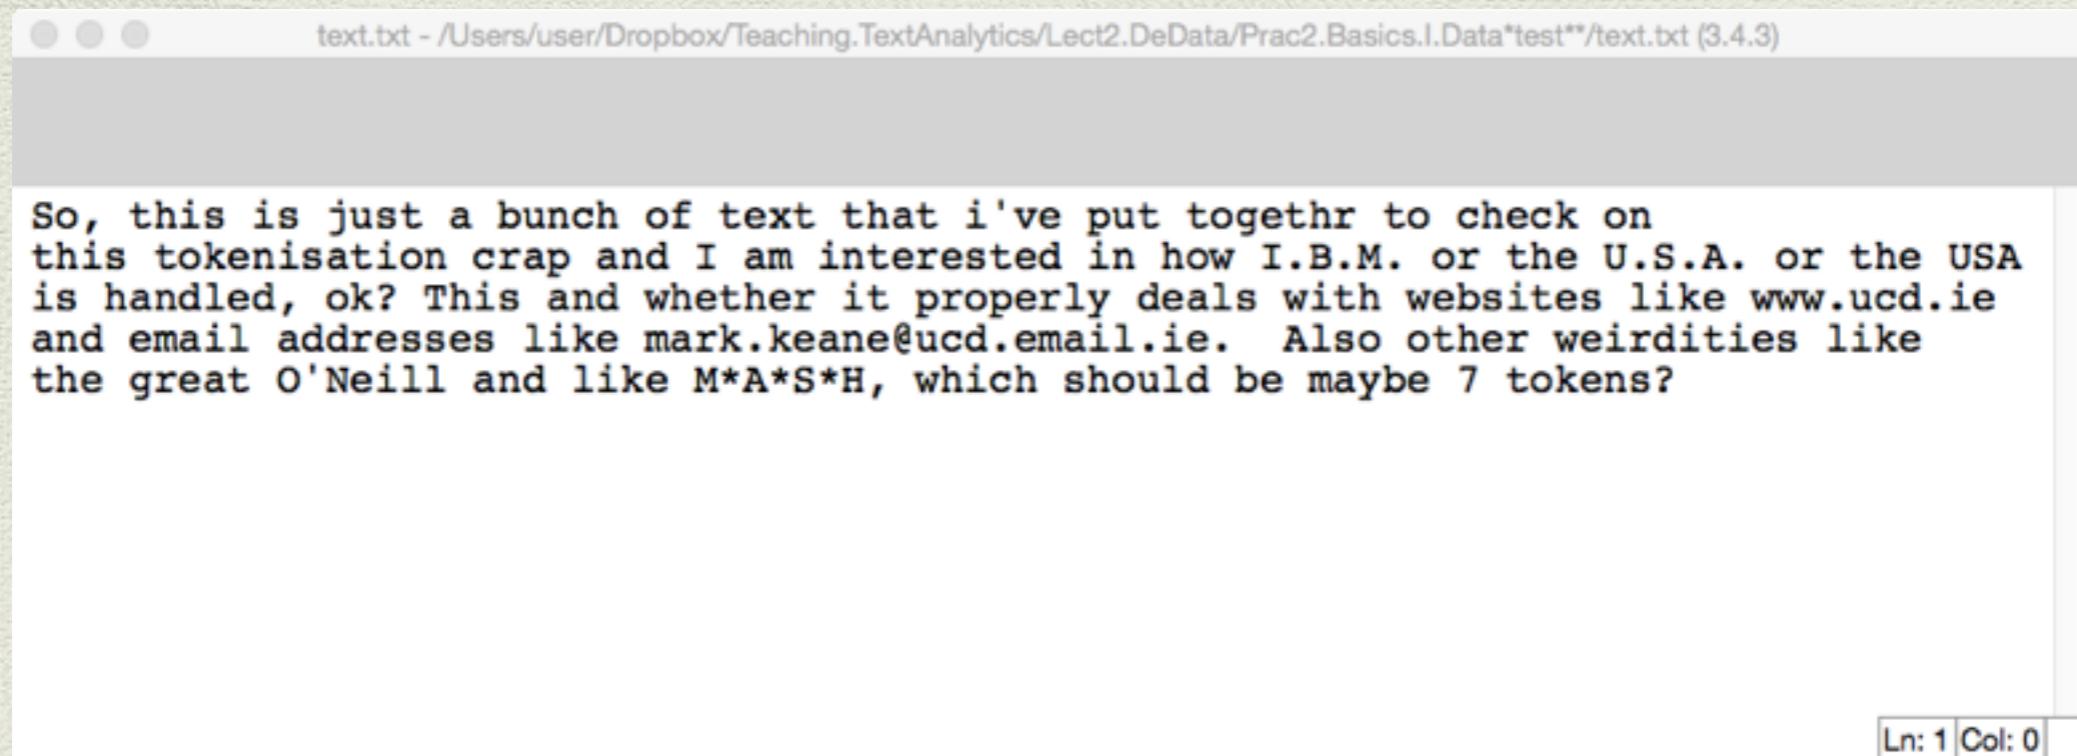
www.ucd.ie: 'www' 'ucd' 'ie' / 'www.ucd.ie' / 'www' 'ucd.ie'

mark.keane@ucd.ie: 'mark' 'keane' '@' 'ucd' 'ie' / what?

- ◆ Which one you pick has impact on the text data

Tokenising I

- ◆ Tokenisers work off general rules and a lot of really quite specific ones too
- ◆ Consider the following text file:

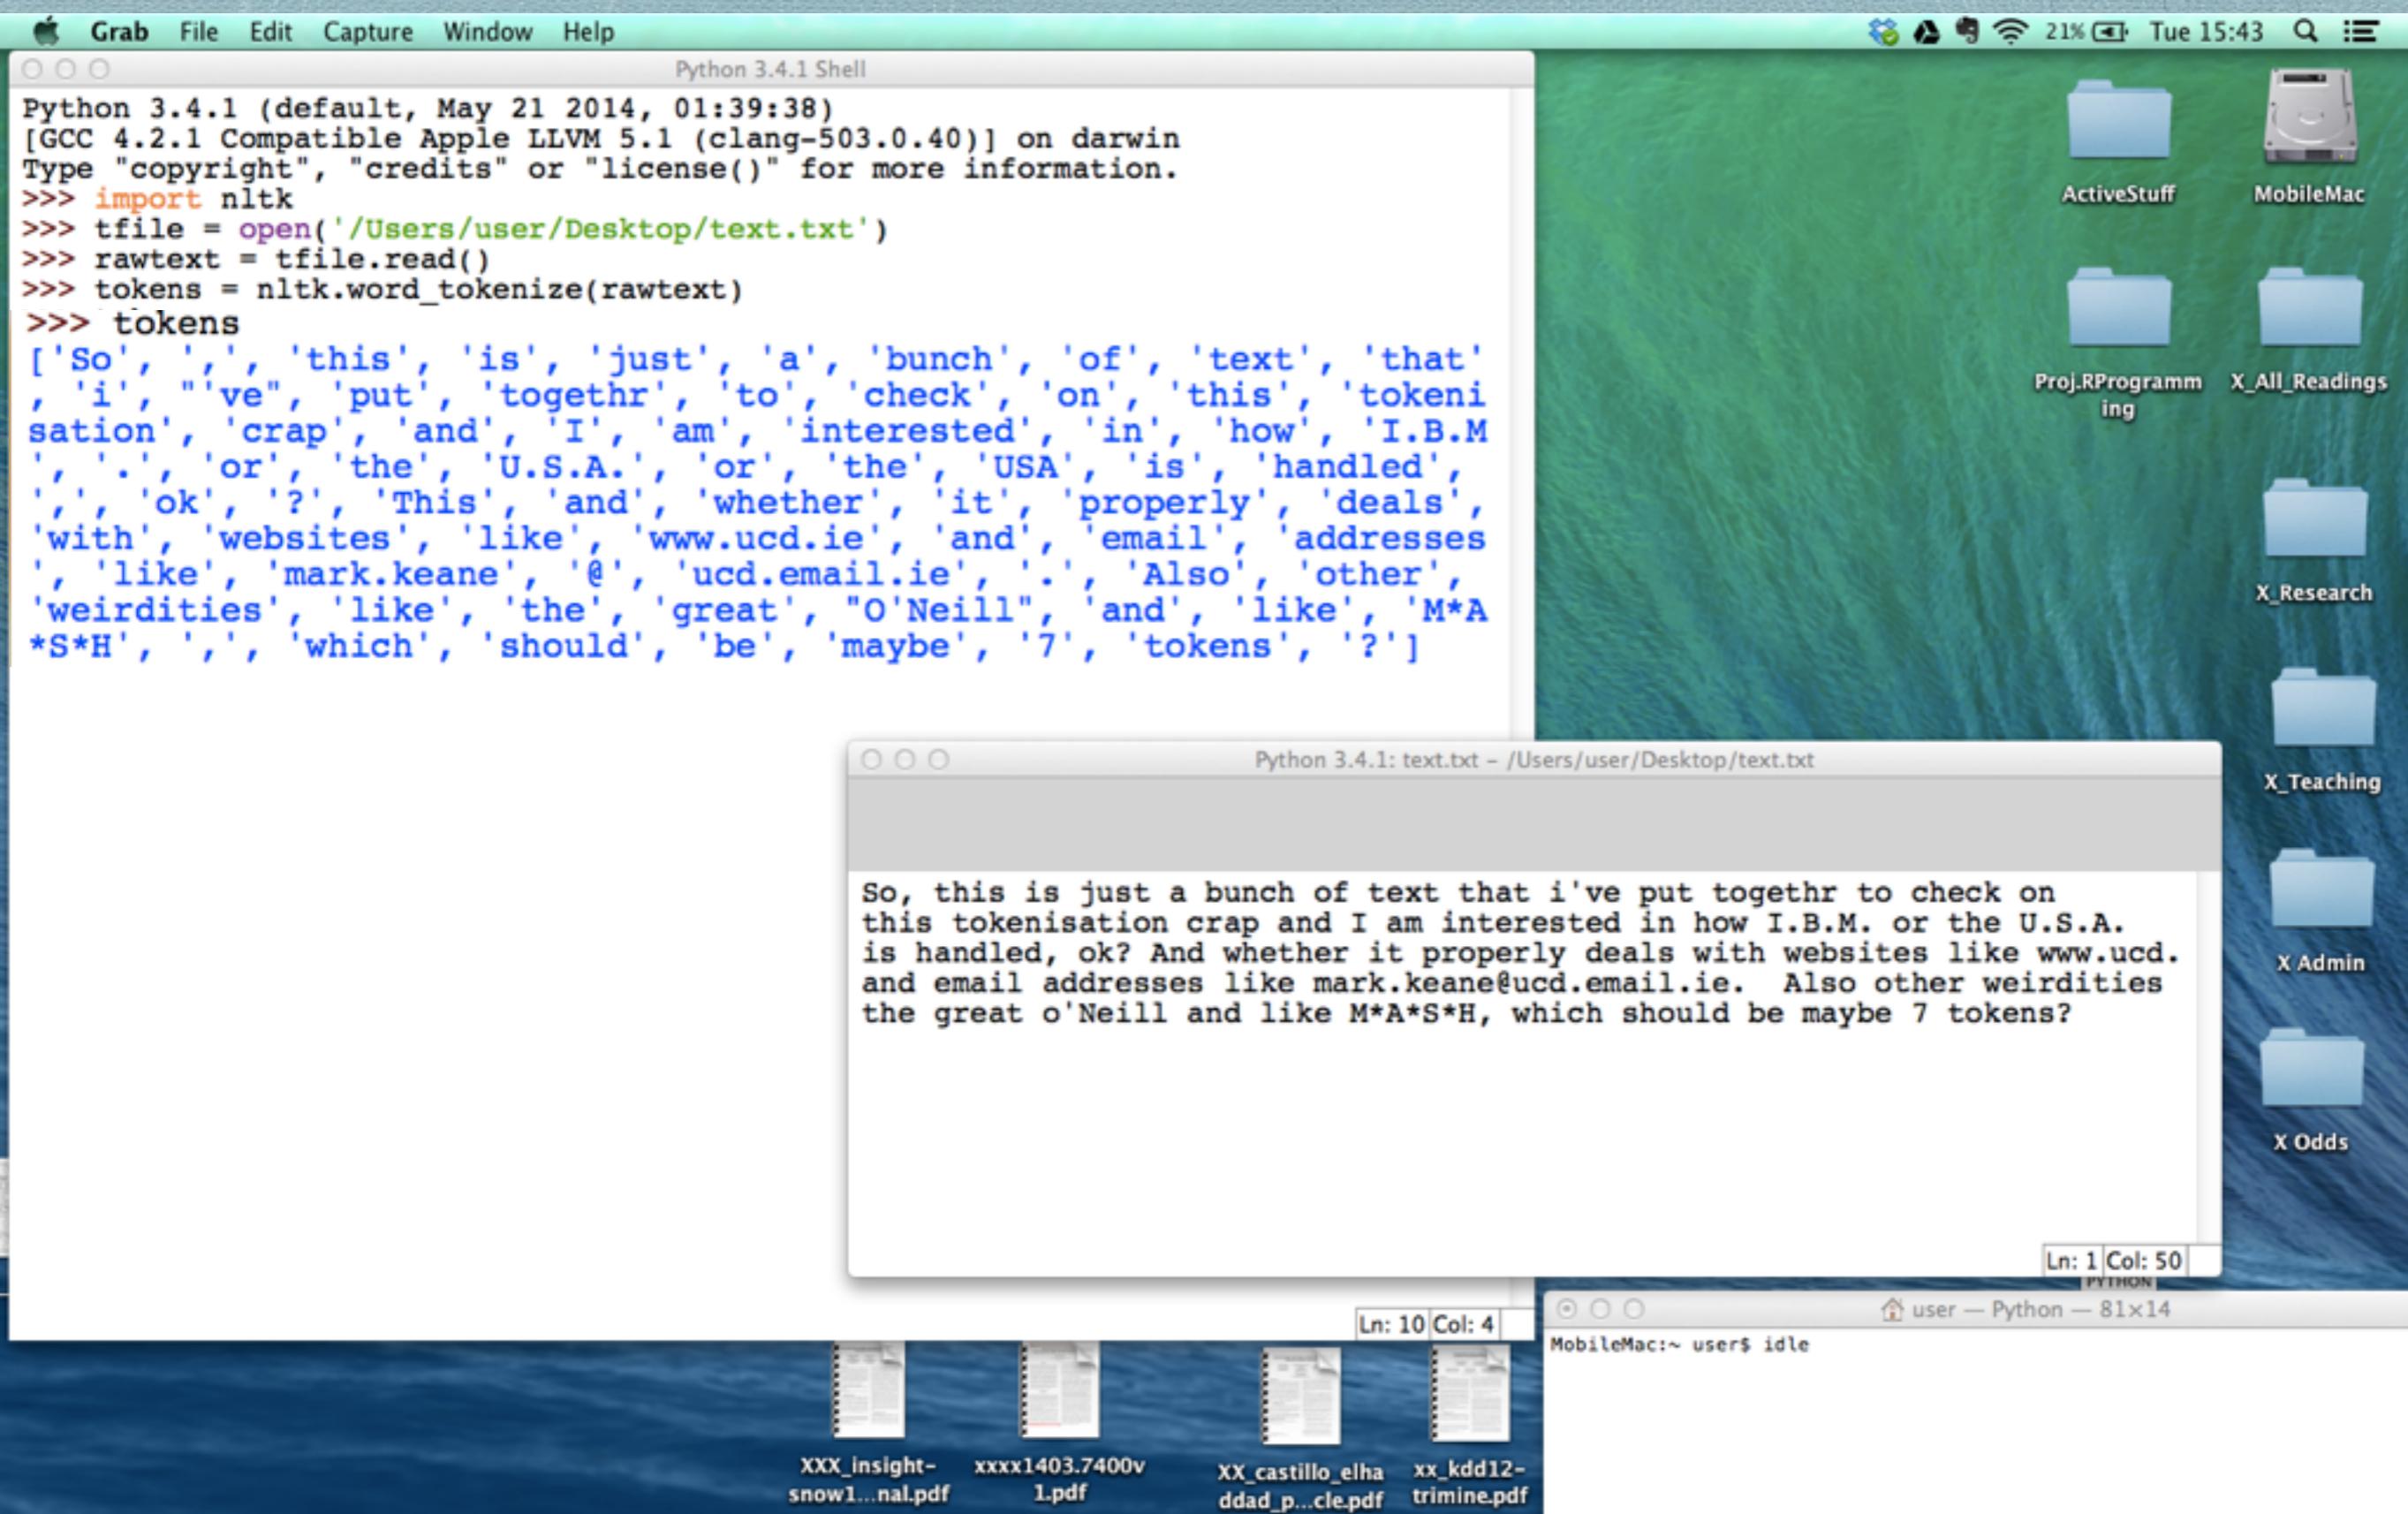


The screenshot shows a text editor window with the following details:

- File path: text.txt - /Users/user/Dropbox/Teaching.TextAnalytics/Lect2.DeData/Prac2.Basics.I.Data/*test*/text.txt (3.4.3)
- Text content:

```
So, this is just a bunch of text that i've put together to check on
this tokenisation crap and I am interested in how I.B.M. or the U.S.A. or the USA
is handled, ok? This and whether it properly deals with websites like www.ucd.ie
and email addresses like mark.keane@ucd.email.ie. Also other weirdities like
the great O'Neill and like M*A*S*H, which should be maybe 7 tokens?
```
- Status bar: Ln: 1 Col: 0

Tokenising III



Tokenising III

- ◆ So, the **nltk** Python tokeniser obviously works off a set of rules for handling things of different types
- ◆ Some rules are quite general (eg handling fullstops)
- ◆ Other rules are very specific ($M^*A^*S^*H$)
- ◆ **nltk** makes it easy for you to extend these rules; using the **re** package (for regular expressions)

Some Downsides...

- ◆ There are other problems to handle in English, like hyphenation, foreign phrases (*au fait*), numbers, acronyms, etc...
- ◆ Problem: solution is quite language-dependent, may need different solutions in:
 - ◆ German: where they compound a lot
 - ◆ Chinese/Korean: where clear word boundaries are not shown by spaces

Why Normalise ?

- ◆ Is often done to reduce of minor differences between tokens, like:
 - ◆ accented differences (cliche, cliché)
 - ◆ capitals (the, The)
 - ◆ acronyms (U.S.A., USA)
- ◆ Though there may be more complex solutions involving *equivalence classes* WHAT?EG?

nltk Normalisation = downcase

The screenshot shows a Mac desktop environment with a green and blue abstract background. On the left, there's a vertical column of icons for 'Post Adverts', 'Proj.GroundTruth', 'Proj.Python', 'Teaching.TextAnalytics', and '_Flow.gif'. On the right, there's a file browser window titled 'sd' showing various folders like 'ActiveStuff', 'MobileMac', 'Proj.RProgramming', 'X_All_Readings', 'tokens', 'X_Research', 'file', 'X_Teaching', 'text.txt', 'X_Admin', 'ryanneDoyle', 'X_Odds', and 'PYTHON'. In the center, there are two terminal windows. The top window is titled 'Python 3.4.1 Shell' and contains the following code and output:

```
"So, this is just a bunch of text that i've put togethr to check on\nthis tokenisation crap and I am interested in how I.B.M. or the U.S.A. or the USA\nis handled, ok? This and whether it properly deals with websites like www.ucd.ie\nand email addresses like mark.keane@ucd.email.ie. Also other weirdities like\nthe great O'Neill and like M*A*S*H, which should be maybe 7 tokens?\n"
>>> tokens = nltk.word_tokenize(rawtext)
>>> tokens
['So', ' ', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', 'i', "'ve",
 'put', 'togethr', 'to', 'check', 'on', 'this', 'tokenisation', 'crap', 'and', 'I',
 'am', 'interested', 'in', 'how', 'I.B.M.', ' ', 'or', 'the', 'U.S.A.', 'or', 'th',
 'e', 'USA', 'is', 'handled', ' ', 'ok', '?', 'This', 'and', 'whether', 'it', 'prop',
 'erly', 'deals', 'with', 'websites', 'like', 'www.ucd.ie', 'and', 'email', 'addres',
 'ses', 'like', 'mark.keane', '@', 'ucd.email.ie', ' ', 'Also', 'other', 'weirditie',
 's', 'like', 'the', 'great', "O'Neill", 'and', 'like', 'M*A*S*H', ' ', 'which', 's',
 'hould', 'be', 'maybe', '7', 'tokens', '?']
>>> words = [w.lower() for w in tokens]
>>> words
['so', ' ', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', 'i', "'ve",
 'put', 'togethr', 'to', 'check', 'on', 'this', 'tokenisation', 'crap', 'and', 'i',
 'am', 'interested', 'in', 'how', 'i.b.m.', ' ', 'or', 'the', 'u.s.a.', 'or', 'th',
 'e', 'usa', 'is', 'handled', ' ', 'ok', '?', 'this', 'and', 'whether', 'it', 'prop',
 'erly', 'deals', 'with', 'websites', 'like', 'www.ucd.ie', 'and', 'email', 'addres',
 'ses', 'like', 'mark.keane', '@', 'ucd.email.ie', ' ', 'also', 'other', 'weirditie',
 's', 'like', 'the', 'great', "o'nellig", 'and', 'like', 'm*a*s*h', ' ', 'which', 's',
 'hould', 'be', 'maybe', '7', 'tokens', '?']
```

The bottom window is titled 'Python 3.4.1: text.txt - /Users/user/Desktop/text.txt' and contains the text:

```
So, this is just a bunch of text that i've put togethr to check on
this tokenisation crap and I am interested in how I.B.M. or the U.S.A. or the USA
is handled, ok? This and whether it properly deals with websites like www.ucd.ie
and email addresses like mark.keane@ucd.email.ie. Also other weirdities like
the great O'Neill and like M*A*S*H, which should be maybe 7 tokens?
```

At the bottom of the screen, there are several status icons and a small text area at the bottom right.

Text Pre-Processing

Misspelling

Why Fix Spellings ?

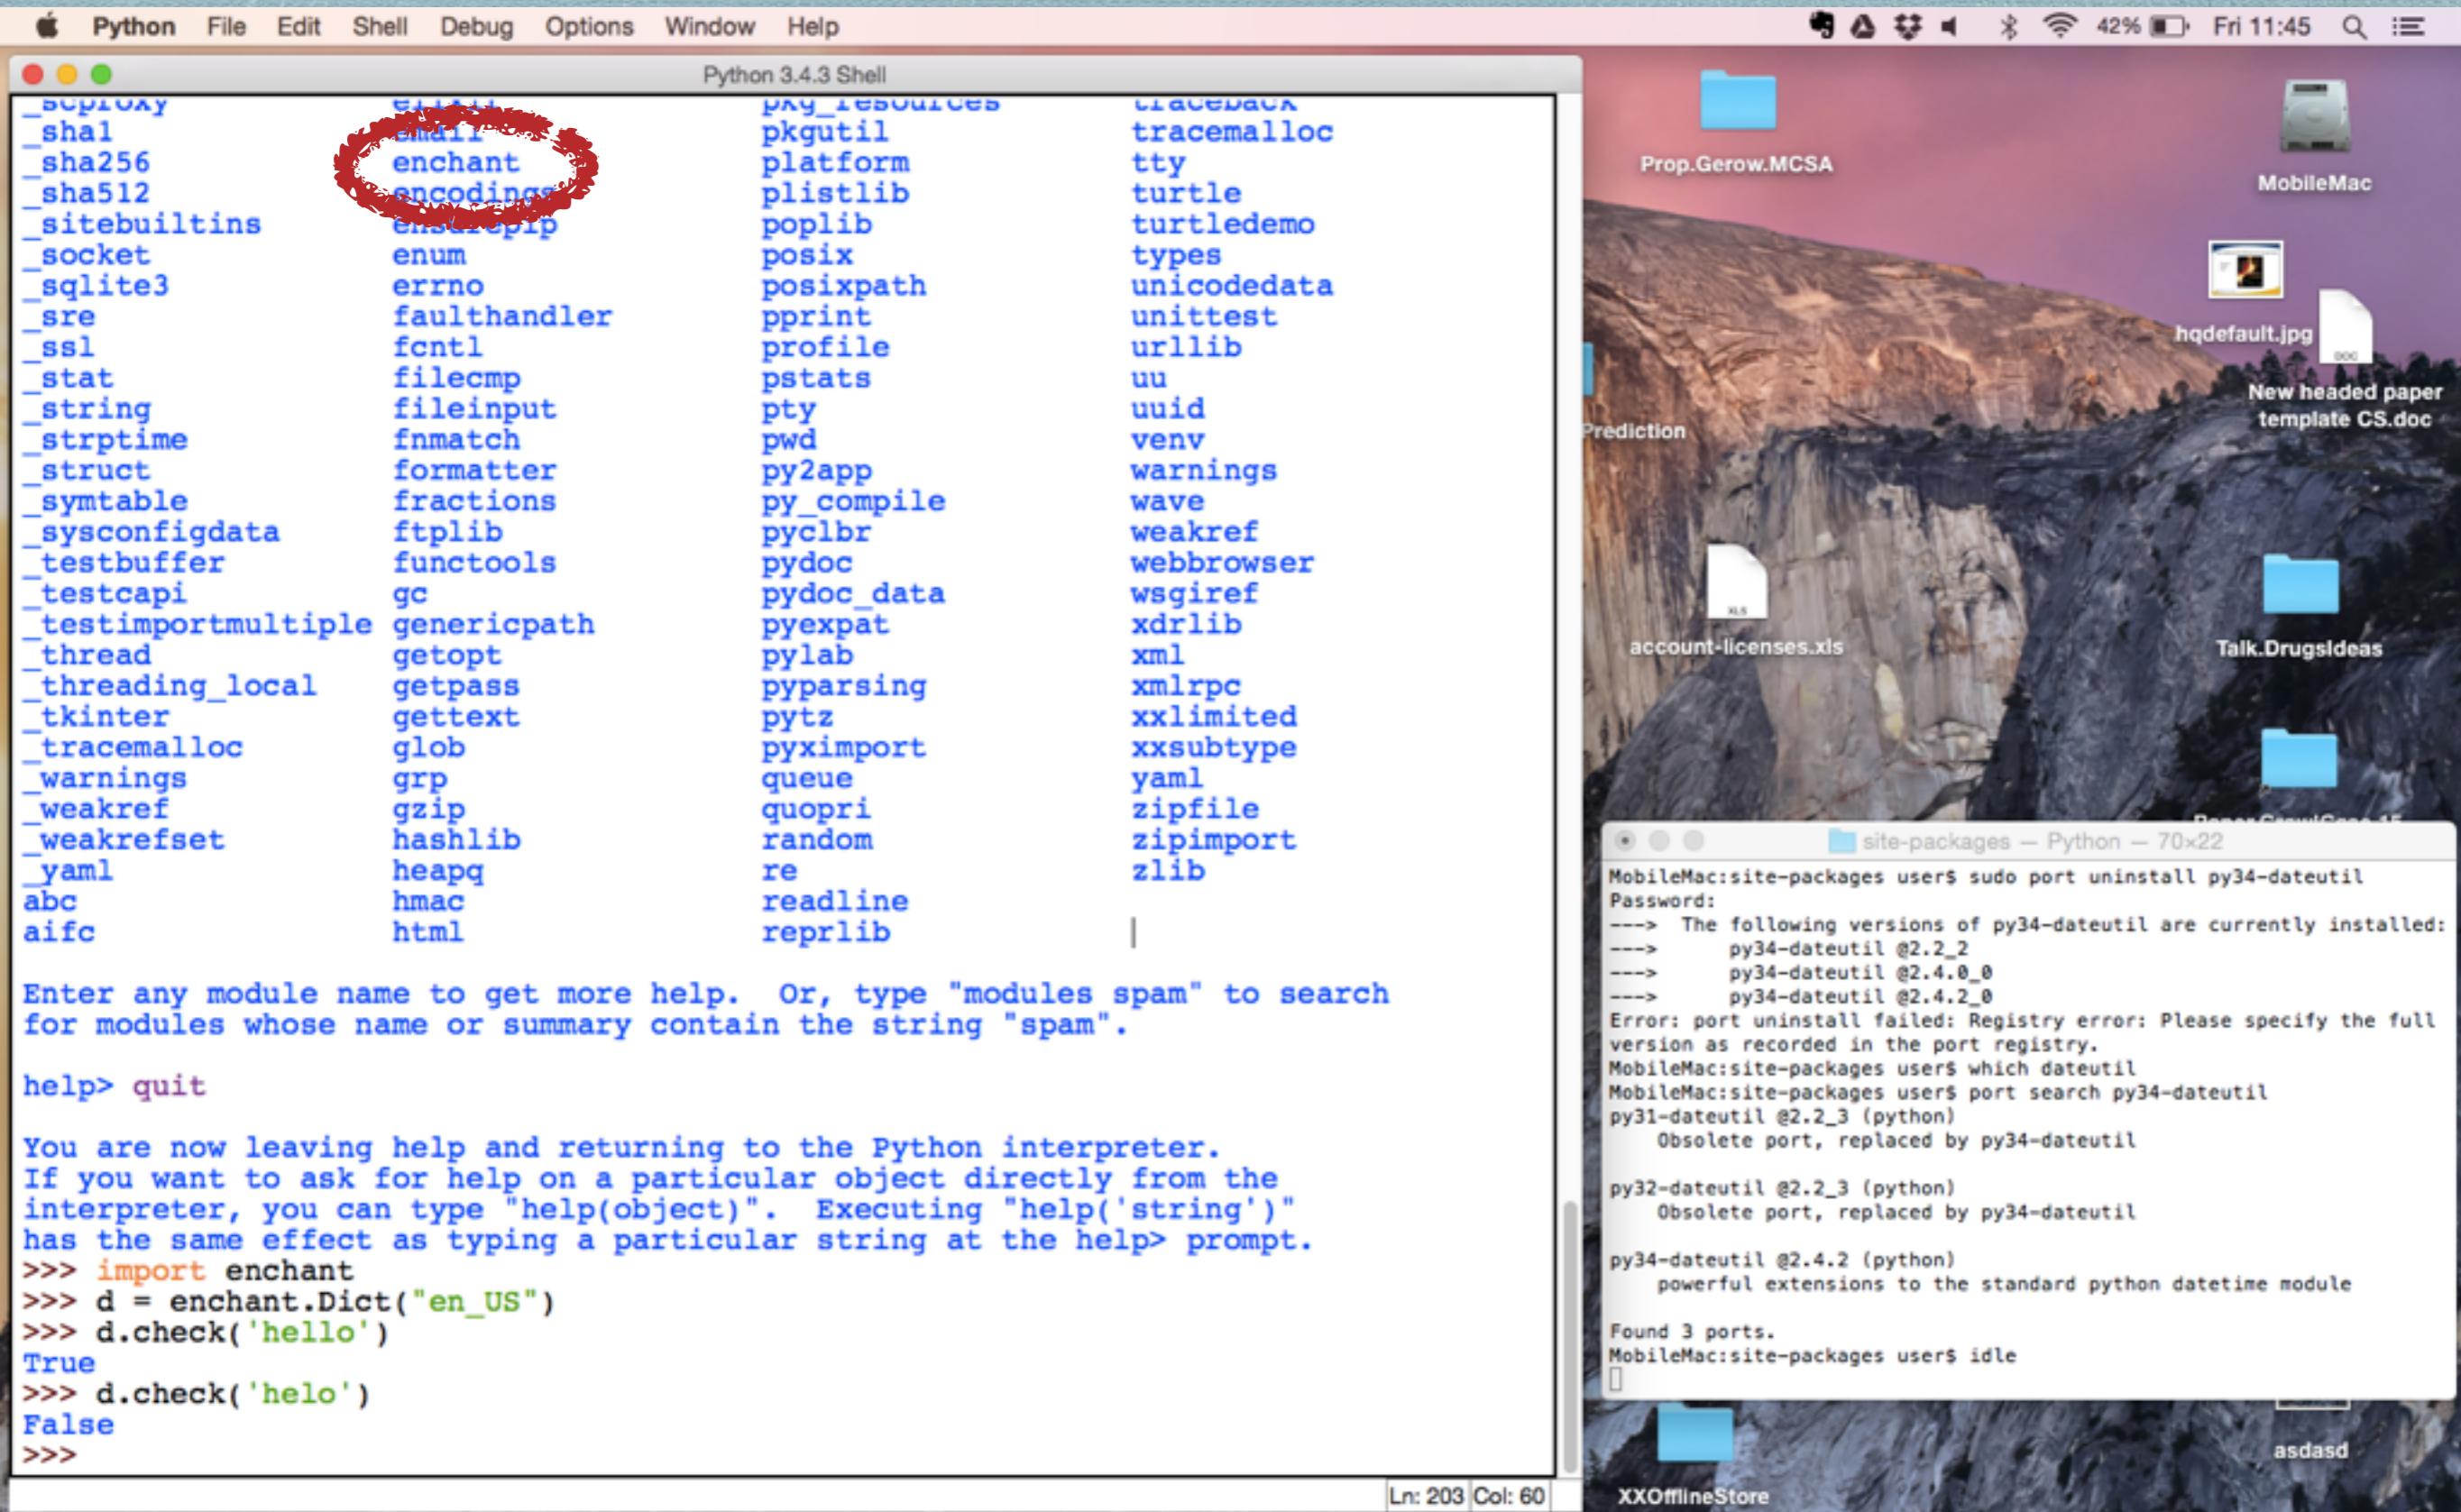
- ◆ Obviously, we cannot assume all the words are spelt correctly; books best, tweets worst
- ◆ So, after tokenizing and normalisation, one could run a spelling checker over the tokens
- ◆ Standard solution is to take a dictionary, compare each word against it, if no entry found, use shortest edit distance to a word in dict

How to Fix Spellings

- ◆ Python has a number of different packages that do spelling correction and front-ends that will highlight misspellings in text
- ◆ the **enchant** package is one of most popular
- ◆ So, PC/Mac Install Enchant

```
$ sudo port install y34-enchant
```

Spellings: enchant module



Python 3.4.3 Shell

```
_scpi
_sha1
_sha256
_sha512
_sitebuiltins
_socket
_sqlite3
_sre
_ssl
_stat
_string
_strptime
_struct
_symtable
_sysconfigdata
_testbuffer
_testcapi
_testimportmultiple
_thread
_threading_local
_tkinter
_tracemalloc
_warnings
_weakref
_weakrefset
_yaml
abc
aifc
_environ
_md5
_enchant
_encodings
_enstdepip
enum
errno
faulthandler
fcntl
filecmp
fileinput
fnmatch
formatter
fractions
ftplib
functools
gc
genericpath
 getopt
getpass
gettext
glob
grp
gzip
hashlib
heapq
hmac
html
_environ
_md5
_enchant
_encodings
_enstdepip
platform
plistlib
poplib
posix
posixpath
pprint
profile
pstats
pty
pwd
py2app
py_compile
pyclbr
pydoc
pydoc_data
pyexpat
pylab
pyparsing
pytz
pyximport
queue
quopri
random
re
readline
reprlib
tracedash
tracemalloc
tty
turtle
turtledemo
types
unicodedata
unittest
urlib
uu
uuid
venv
warnings
wave
weakref
webbrowser
wsgiref
xdrlib
xml
xmlrpclib
xxlimited
xxsubtype
yaml
zipfile
zipimport
zlib

Enter any module name to get more help. Or, type "modules spam" to search
for modules whose name or summary contain the string "spam".

help> quit

You are now leaving help and returning to the Python interpreter.
If you want to ask for help on a particular object directly from the
interpreter, you can type "help(object)". Executing "help('string')"
has the same effect as typing a particular string at the help> prompt.

>>> import enchant
>>> d = enchant.Dict("en_US")
>>> d.check('hello')
True
>>> d.check('heloo')
False
>>>
```

Prop.Gerow.MCSA

MobileMac

hqdefault.jpg

New headed paper template CS.doc

Talk.Drugsideas

account-licenses.xls

site-packages — Python — 70x22

```
MobileMac:site-packages user$ sudo port uninstall py34-dateutil
Password:
--> The following versions of py34-dateutil are currently installed:
--> py34-dateutil @2.2_2
--> py34-dateutil @2.4.0_0
--> py34-dateutil @2.4.2_0
Error: port uninstall failed: Registry error: Please specify the full
version as recorded in the port registry.
MobileMac:site-packages user$ which dateutil
MobileMac:site-packages user$ port search py34-dateutil
py31-dateutil @2.2_3 (python)
    Obsolete port, replaced by py34-dateutil
py32-dateutil @2.2_3 (python)
    Obsolete port, replaced by py34-dateutil
py34-dateutil @2.4.2 (python)
    powerful extensions to the standard python datetime module

Found 3 ports.
MobileMac:site-packages user$ idle
```

XXOfflineStore

asdasd

Fixing Spellings: BigData Way

The image shows two side-by-side Python shells running on a Mac OS X desktop. The left shell is titled 'Python 3.4.1 Shell' and contains a session of testing a 'correct' function with misspelled words like 'spellingd', 'stroeked', and 'streked'. The right shell is titled 'Python 3.4.1: norvigspell.py - /Users/user/Desktop/norvigspell.py' and displays the source code for the Norvig Spelling Corrector. The code defines functions for training a model on a corpus of words, generating edits for a given word, and finding the most likely correct spelling based on edit distances.

```
Python 3.4.1 (default, May 21 2014, 01:39:38)
[GCC 4.2.1 Compatible Apple LLVM 5.1 (clang-503.0.40
)] on darwin
Type "copyright", "credits" or "license()" for more
information.

>>> ===== RESTART =====
>>>
>>> correct('spellingd')
'spelling'
>>> correct('stroeked')
'stroked'
>>> correct('streked')
'stroked'
>>> correct('stoeked')
'stocked'
>>>

# Norvig Spelling Corrector

import re, collections

def words(text): return re.findall('[a-z]+', text.lower())

def train(features):
    model = collections.defaultdict(lambda: 1)
    for f in features:
        model[f] += 1
    return model

NWORDS = train(words(open('/Users/user/Desktop/big.txt').read()))

alphabet = 'abcdefghijklmnopqrstuvwxyz'

def edits1(word):
    splits      = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes     = [a + b[1:] for a, b in splits if b]
    transposes = [a + b[1] + b[0] + b[2:] for a, b in splits if len(b)>1]
    replaces   = [a + c + b[1:] for a, b in splits for c in alphabet if b]
    inserts    = [a + c + b for a, b in splits for c in alphabet]
    return set(deletes + transposes + replaces + inserts)

def known_edits2(word):
    return set(e2 for e1 in edits1(word) for e2 in edits1(e1) if e2 in NWORDS)

def known(words): return set(w for w in words if w in NWORDS)

def correct(word):
    candidates = known([word]) or known(edits1(word)) or known_edits2(word)
    return max(candidates, key=NWORDS.get)
```

```
MobileMac:~ users$ idle
MobileMac:~ users$ python27
-bash: python27: command not found
MobileMac:~ users$ python27
-bash: python27: command not found
MobileMac:~ users$ idle
```

How It Works: Some Probability Theory

How does it work? First, a little theory. Given a word, we are trying to choose the most likely spelling correction for that word (the "correction" may be the original word itself). There is no way to know for sure (for example, should "lates" be corrected to "late" or "latest"?), which suggests we use probabilities. We will say that we are trying to find the correction c , out of all possible corrections, that maximizes the probability of c given the original word w :

$$\operatorname{argmax}_c P(c|w)$$

By [Bayes' Theorem](#) this is equivalent to:

$$\operatorname{argmax}_c P(w|c) P(c) / P(w)$$

Since $P(w)$ is the same for every possible c , we can ignore it, giving:

$$\operatorname{argmax}_c P(w|c) P(c)$$

There are three parts of this expression. From right to left, we have:

1. $P(c)$, the probability that a proposed correction c stands on its own. This is called the **language model**: think of it as answering the question "how likely is c to appear in an English text?" So $P("the")$ would have a relatively high probability, while $P("zxzxzxzyy")$ would be near zero.
2. $P(w|c)$, the probability that w would be typed in a text when the author meant c . This is the **error model**: think of it as answering "how likely is it that the author would type w by mistake when c was intended?"
3. argmax_c , the control mechanism, which says to enumerate all feasible values of c , and then choose the one that gives the best combined probability score.

One obvious question is: why take a simple expression like $P(c|w)$ and replace it with a more complex expression involving two models rather than one? The answer is that $P(c|w)$ is *already* conflating two factors, and it is easier to separate the two out and deal with them explicitly. Consider the misspelled word $w="thew"$ and the two candidate corrections $c="the"$ and $c="thaw"$. Which has a higher $P(c|w)$? Well, "thaw" seems good because the only change is "a" to "e", which is a small change. On the other hand, "the" seems good because "the" is a very common word, and perhaps the typist's finger slipped off the "e" onto the "w". The point is that to estimate $P(c|w)$ we have to consider both the probability of c and the probability of the change from c to w anyway, so it is cleaner to formally separate the two factors.

Now we are ready to show how the program works. First $P(c)$. We will read a big text file, [big.txt](#), which consists of about a million words. The file is a concatenation of several public domain books from [Project Gutenberg](#) and lists of most frequent words from [Wiktionary](#) and the [British National Corpus](#). (On the plane home when I was writing the first version of the code all I had was a collection of Sherlock Holmes stories that happened to be on my laptop; I added the other sources later and stopped adding texts when they stopped helping, as we shall see in the Evaluation section.)

We then extract the individual words from the file (using the function `words`, which converts everything to lowercase, so that "the" and "The" will be the same and then defines a word as a sequence of alphabetic characters, so "don't" will be seen as the two words "don" and "t"). Next we train a probability model, which is a fancy way of saying we count how many times each word occurs, using the function `train`. It looks like this:

```
def words(text): return re.findall('[a-z]+', text.lower())

def train(features):
    model = collections.defaultdict(lambda: 1)
    for f in features:
        model[f] += 1
    return model

NWORDS = train(words(file('big.txt').read()))
```

At this point, `NWORDS[w]` holds a count of how many times the word `w` has been seen. There is one complication: novel words. What happens with a perfectly good word of English that wasn't seen in our training data? It would be bad form to say the probability of a word is zero just because we haven't seen it yet. There are several standard approaches to this problem; we take the easiest one, which is to treat novel words as if we had seen them once. This general process is called **smoothing**, because we are smoothing over the parts of the probability distribution that would have been zero, bumping them up to the smallest possible count. This is achieved through the class `collections.defaultdict`, which is like a regular Python dict (what other languages call hash tables) except that we can specify the default value of any key; here we use 1.

Now let's look at the problem of enumerating the possible corrections c of a given word w . It is common to talk of the **edit distance** between two words: the number of edits it would take to turn one into the other. An edit can be a deletion (remove one letter), a transposition (swap adjacent letters), an alteration (change one letter to another) or an insertion (add a letter). Here's a function that returns a set of all words c that are one edit away from w :

```
def edits1(word):
    splits      = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes     = [a + b[1:] for a, b in splits if b]
    transposes = [a + b[1] + b[0] + b[2:] for a, b in splits if len(b)>1]
    replaces   = [a + c + b[1:] for a, b in splits for c in alphabet if b]
    inserts    = [a + c + b      for a, b in splits for c in alphabet]
    return set(deletes + transposes + replaces + inserts)
```

Fixing Spellings: Actual Prog

The image shows a Mac OS X desktop with three windows open:

- Python 3.4.3 Shell**: A terminal window showing Python 3.4.3 running on a Mac (GCC 4.2.1 Compatible Apple LLVM 6.0). It displays a spell checker script and some test cases.
- norvigspell.py**: A code editor window containing the Norvig Spelling Corrector script. The code defines functions for training a model, generating edits, and finding known words.
- site-packages**: A terminal window showing the output of `idle` commands, indicating the presence of the `pytz` package.

```
Python 3.4.3 (default, May 25 2015, 18:48:21)
[GCC 4.2.1 Compatible Apple LLVM 6.0 (clang-600.0.56)] on darwin
Type "copyright", "credits" or "license()" for more information.
>>> ===== RESTART =====
>>>
/Users/user/Dropbox/Teaching.TextAnalytics/Lect2.DeData/
Prac2.Basics.I.Data*test**
/Users/user/Dropbox/Teaching.TextAnalytics/Lect2.DeData/
Prac2.Basics.I.Data*test**/big.txt
>>> correct('hello')
'hello'
>>> correct('hel0')
'held'
>>> correct('streked')
'stroked'
>>> correct('stroked')
'stroked'
>>> |
```

```
# Norvig Spelling Corrector
import re, collections, os

__location__ = os.path.realpath(
    os.path.join(os.getcwd(),os.path.dirname(__file__)))
print(__location__)

file_name = os.path.join(__location__, 'big.txt')
print(file_name)

def words(text): return re.findall('[a-z]+', text.lower())

def train(features):
    model = collections.defaultdict(lambda: 1)
    for f in features:
        model[f] += 1
    return model

NWORDS = train(words(open(file_name).read()))

alphabet = 'abcdefghijklmnopqrstuvwxyz'

def edits1(word):
    splits      = [(word[:i], word[i:]) for i in range(len(word) + 1)]
    deletes     = [a + b[1:] for a, b in splits if b]
    transposes = [a + b[1] + b[0] + b[2:] for a, b in splits if len(b)>1]
    replaces   = [a + c + b[1:] for a, b in splits for c in alphabet if b]
    inserts    = [a + c + b      for a, b in splits for c in alphabet]
    return set(deletes + transposes + replaces + inserts)

def known_edits2(word):
    return set(e2 for e1 in edits1(word) for e2 in edits1(e1) if e2 in NWORDS)

def known(words): return set(w for w in words if w in NWORDS)
```

A terminal window titled "site-packages - Python - 162x6" is visible at the bottom of the screen. It shows the output of the command `idle`, which lists three ports found.

```
powerful extensions to the standard python datetime module
Found 3 ports.
MobileMac:site-packages user$ idle
MobileMac:site-packages user$ idle
```

Fixing Spellings: Google Way

- ◆ [http://www.internetmarketingninjas.com/
blog/search-engine-optimization/google-
spell-check/](http://www.internetmarketingninjas.com/blog/search-engine-optimization/google-spell-check/)

Text Pre-Processing

Stemming & Lemmatising

What is the Problem?

- ◆ We have a number of normalised tokens that, we hope, are the words in the document
- ◆ But, we are going to be counting them and doing other manipulations on frequency
- ◆ So, it is very important that words that are more or less the same are identified...

Why Stem?

- ◆ So, we need to recognise variations of the same *stem* or *root**: fish~ for fishes, fishing, fished...
- ◆ We also need to recognise when two words are the same but from different syntactic categories (or parts of speech, POS); fish the *noun* and fish the *verb*
- ◆ Note, the root form of a word may be quite different **be~** for is, are, am

* may be used differently

Why Stem?

- ◆ So, if we convert the variations into the root form then we have a more accurate count of the word
- ◆ So, if we distinguish two words that appear the same but are actually different fish-v and fish-n (stemming doesn't really do this...)
- ◆ Generally, stemming refers to the recovery of the root forms of words to show these commonalities and differences

Stemming Definition...

Stemming

From Wikipedia, the free encyclopedia



This article **needs attention from an expert on the subject**.

Please add a *reason* or a *talk* parameter to this template to explain the issue with the article. Consider [associating this request](#) with a [WikiProject](#). (October 2010)

For the skiing technique, see [Stem \(skiing\)](#).

In [linguistic morphology](#) and [information retrieval](#), **stemming** is the process for reducing inflected (or sometimes derived) words to their stem, base or root form—generally a written word form. The stem need not be identical to the [morphological root](#) of the word; it is usually sufficient that related words map to the same stem, even if this stem is not in itself a valid root. [Algorithms](#) for stemming have been studied in [computer science](#) since the 1960s. Many [search engines](#) treat words with the same stem as [synonyms](#) as a kind of query expansion, a process called **conflation**.

Stemming programs are commonly referred to as **stemming algorithms** or **stemmers**.

Stemming Algorithms

- ◆ Technically, stemming is described the removal of morphological affixes
- ◆ Porter Stemmer, Snowball stemmers, most used but many more in **nltk** (types and languages)
- ◆ [http://www.nltk.org/
howto/stem.html](http://www.nltk.org/howto/stem.html)

affix

verb

3rd person present: **affixes**
/ə'fɪks/

1. stick, attach, or fasten (something) to something else.
"panels to which he affixes copies of fine old prints"
synonyms: [attach](#), [stick](#), [fasten](#), [bind](#), [fix](#), [post](#), [secure](#), [join](#), [connect](#), [couple](#);
[More](#)

noun GRAMMAR

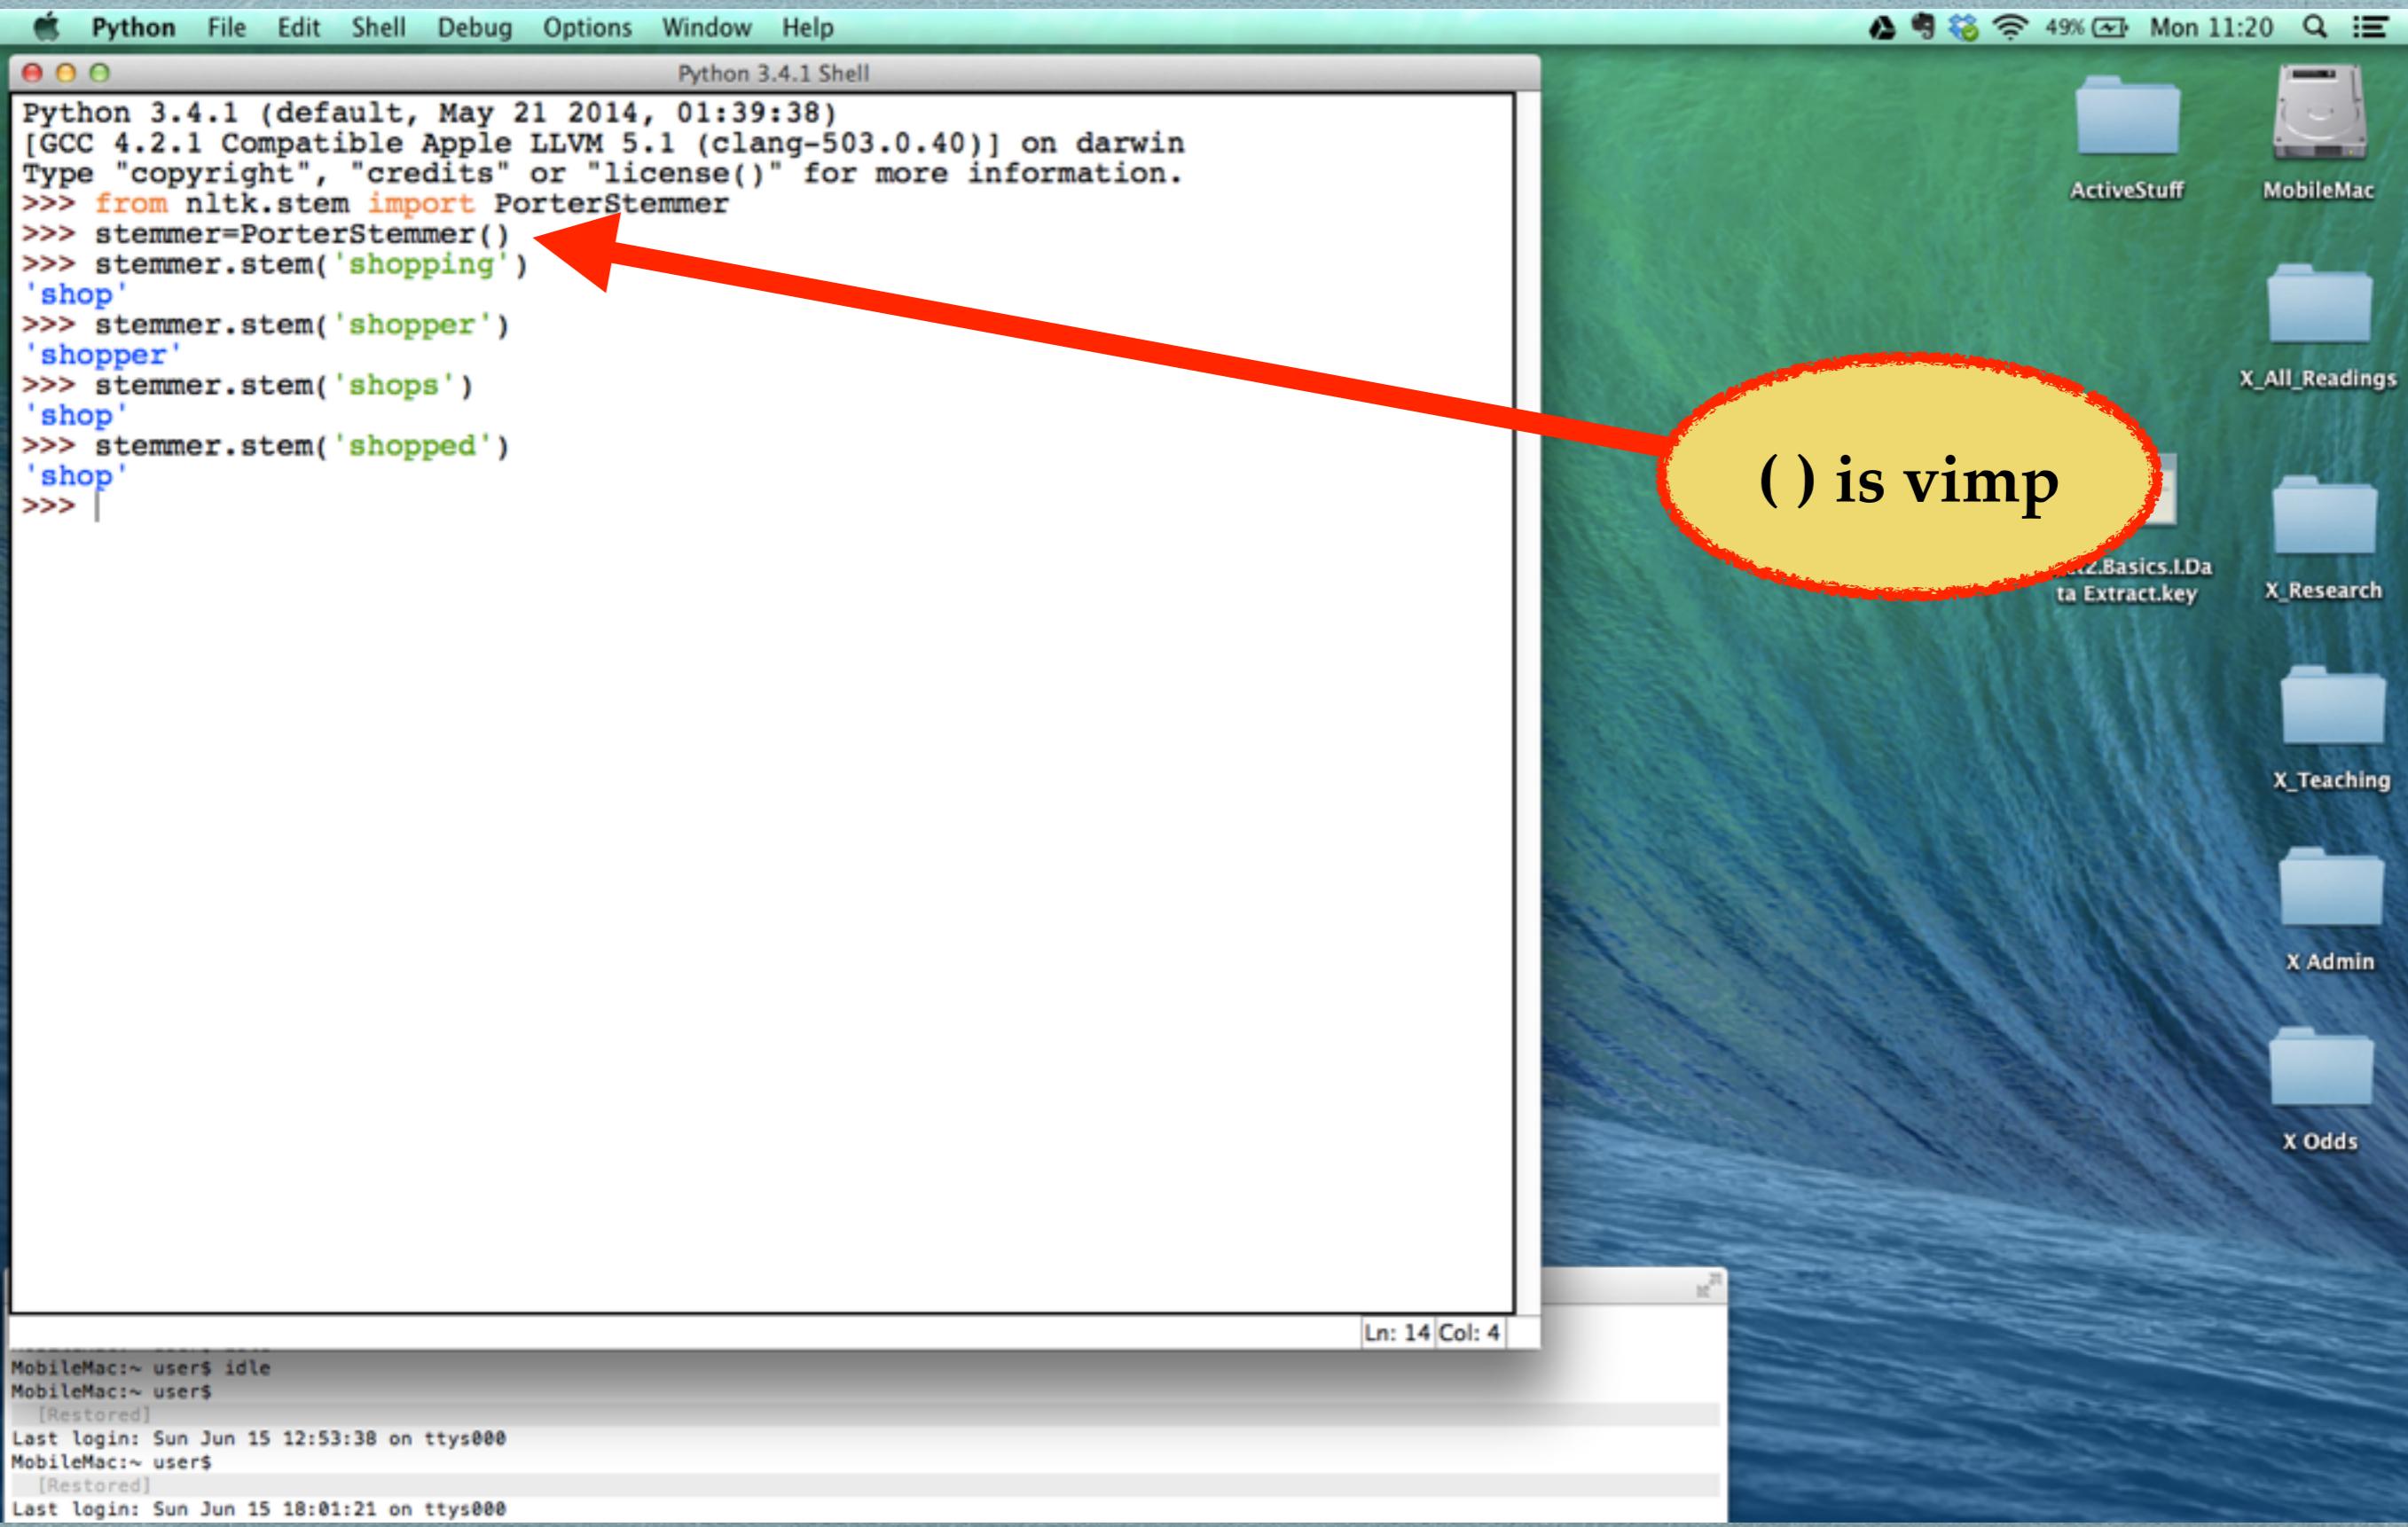
plural noun: **affixes**
/'afɪks/

1. an addition to the base form or stem of a word in order to modify its meaning or create a new word.

Porter Stemming: Brutal

- ◆ Porter Stemming is quite brutal, tho' most of it makes sense (e.g., does plurals well)
- ◆ It works off a mixture of specific rules and very general ones (e.g., cutting '-ed' off words)
- ◆ <http://tartarus.org/martin/PorterStemmer/>

Porter Stemming: Eg



Porter Stemming: Eg

The image shows a Mac OS X desktop environment with several windows open:

- Python 3.4.1 Shell:** A window titled "Python 3.4.1 Shell" containing Python code demonstrating Porter Stemming. The code reads a text file from the desktop, tokenizes it, and then stems each word using the PorterStemmer class from the nltk.stem module.
- Terminal:** A window showing the command-line history of the user. It includes the path "/Users/user/Desktop/text.txt" and the content of the file, which is a bunch of text about stemming.
- File Browser:** A window titled "/Users/user/Desktop/text.txt" showing the file's contents in a text editor.
- Desktop:** The main desktop area showing various icons for folders like "ActiveStuff", "MobileMac", "stem", "X_All_Readings", "Lect2.Basics.IData Extract.key", "X_Research", "X_Teaching", and "X Admin".

Porter Stemming: Eg

```
Python 3.4.1 (default, May 21 2014, 01:39:38)
[GCC 4.2.1 Compatible Apple LLVM 5.1 (clang-503.0.40)] on darwin
Type "copyright", "credits" or "license()" for more information.
>>> import nltk
>>> text = open('/Users/user/Desktop/text.txt')
>>> rawtext = text.read()
>>> from nltk.stem import PorterStemmer
>>> rawtext
"So, this is just a bunch of text that i've put together to check on\nthis tokenisation crap and I am interested in how I.B.M. or the U.S.A. or the USA\nis handled, ok? This and whether it properly deals with websites like www.ucd.ie\\nand email addresses like mark.keane@ucd.email.ie. Also other weirdities like\\nthe great O'Neill and like M*A*S*H, which should be maybe 7 tokens?\n"
>>> splitrawtext = rawtext.split(" ")
>>> splitrawtext
['So,', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put', 'together', 'to', 'check', 'on\\nthis', 'tokenisation', 'crap', 'and', 'I', 'am', 'interested', 'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\\nis', 'handled,', 'ok?', 'This', 'and', 'whether', 'it', 'properly', 'deals', 'with', 'websites', 'like', 'www.ucd.ie\\nand', 'email', 'addresses', 'like', 'mark.keane@ucd.email.ie.', 'Also', 'other', 'weirdities', 'like\\nthe', 'great', "O'Neill", 'and', 'like', 'M*A*S*H,', 'which', 'should', 'be', 'maybe', '7', 'tokens?\n']
>>> stemmer = PorterStemmer()
>>> stemmer.stem(wd) for wd in splitrawtext
SyntaxError: invalid syntax
>>> [stemmer.stem(wd) for wd in splitrawtext]
['So', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i'v", 'put', 't ogether', 'to', 'check', 'on\\n thi', 'tokenis', 'crap', 'and', 'I', 'am', 'interest ', 'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\\ni', 'handled', 'ok?', 'Thi', 'and', 'whether', 'it', 'properli', 'deal', 'with', 'websit', 'l ike', 'www.ucd.ie\\nand', 'email', 'address', 'like', 'mark.keane@ucd.email.ie.', 'Also', 'other', 'weirditi', 'like\\nth', 'great', "O'Neil", 'and', 'like', 'M *A*S*H,', 'which', 'should', 'be', 'mayb', '7', 'tokens?\n']
>>>
```

import the method

Split text into a list of strings

Call stemmer on each string in list

MobileMac:~ user\$ [Restored] Last login: Sun Jun 15 12:53:38 on ttys000 MobileMac:~ user\$ [Restored] Last login: Sun Jun 15 18:01:21 on ttys000 MobileMac:~ user\$ idle

Ln: 17 Col: 53

and email addresses like mark.keane@ucd.email.ie. Also other weirdities like the great O'Neill and like M*A*S*H, which should be maybe 7 tokens?

Porter Stemming: Steps

- ◆ Step 1: Gets rid of plurals and -ed or -ing suffixes
- ◆ Step 2: Turns terminal y to i when there is another vowel in the stem
- ◆ Step 3: Maps double suffixes to single ones:-ization, -ational
- ◆ Step 4: Deals with suffixes, -full, -ness etc.
- ◆ Step 5: Takes off -ant, -ence, etc.
- ◆ Step 6: Removes a final -e

Lancaster Stemming: Eg

The image shows a Mac OS X desktop environment. In the center is a Python 3.4.1 Shell window titled "Python 3.4.1 Shell". The code inside the shell is as follows:

```
>>> splitrawtext
['So,', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put',
'togethr', 'to', 'check', 'on\nthis', 'tokenisation', 'crap', 'and', 'I', 'am',
'interested', 'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\nnis',
'handled', 'ok?', 'This', 'and', 'whether', 'it', 'properly', 'deals', 'with',
'websites', 'like', 'www.ucd.ie\\nand', 'email', 'addresses', 'like', 'mark.keane@ucd.email.ie.',
'', 'Also', 'other', 'weirdities', 'like\\nth', 'great', "O'Neil",
', 'and', 'like', 'M*A*S*H', 'which', 'should', 'be', 'maybe', '7', 'tokens?\n']
>>> stemmer = PorterStemmer()
>>> stemmer.stem(wd) for wd in splitrawtext
SyntaxError: invalid syntax
>>> [stemmer.stem(wd) for wd in splitrawtext]
['So,', 'thi', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i'v", 'put', 't
ogethr', 'to', 'check', 'on\\nthi', 'tokenis', 'crap', 'and', 'I', 'am', 'interest
', 'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\\ni', 'handled,
', 'ok?', 'Thi', 'and', 'whether', 'it', 'properli', 'deal', 'with', 'websit', 'l
ike', 'www.ucd.ie\\nand', 'email', 'address', 'like', 'mark.keane@ucd.email.ie.',
'', 'Also', 'other', 'weirditi', 'like\\nth', 'great', "O'Neil", 'and', 'like', 'M
*A*S*H', 'which', 'should', 'be', 'mayb', '7', 'tokens?\n']
>>> stemmer = LancasterStemmer()
Traceback (most recent call last):
  File "<pyshell#10>", line 1, in <module>
    stemmer = LancasterStemmer()
NameError: name 'LancasterStemmer' is not defined
>>> stemmer = nltk.LancasterStemmer()
>>> [stemmer.stem(wd) for wd in splitrawtext]
['so,', 'thi', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put',
'togethr', 'to', 'check', 'on\\nthis', 'tok', 'crap', 'and', 'i', 'am', 'interest',
'in', 'how', 'i.b.m.', 'or', 'the', 'u.s.a.', 'or', 'the', 'usa\\nis', 'handled',
'ok?', 'thi', 'and', 'wheth', 'it', 'prop', 'deal', 'with', 'websit', 'lik', 'www
.ucd.ie\\nand', 'email', 'address', 'lik', 'mark.keane@ucd.email.ie.', '', 'also',
'oth', 'weird', 'like\\nth', 'gre', "o'neill", 'and', 'lik', 'm*a*s*h', 'which',
'should', 'be', 'mayb', '7', 'tokens?\n']
>>>
>>>
>>>
>>>
>>>
>>>
```

The shell window also shows a command-line interface at the bottom:

```
MobileMac:~ user$ 
[Restored]
Last login: Sun Jun 15 12:53:38 on ttys000
MobileMac:~ user$ 
[Restored]
Last login: Sun Jun 15 18:01:21 on ttys000
MobileMac:~ user$ idle
```

To the right of the shell window is a file browser window showing the desktop. A file named "text.txt" is open in a text editor. The text in the file is:

```
put togethr to check on
in how I.B.M. or the U.S.A. or the USA
ly deals with websites like www.ucd.ie
and email addresses like mark.keane@ucd.email.ie. Also other weirdities like
the great O'Neill and like M*A*S*H, which should be maybe 7 tokens?
```

Lancaster Stemming: Eg

```
Python 3.4.1 Shell
```

```
>>> splitrawtext
['So,', 'this', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put',
'togethr', 'to', 'check', 'on\nthis', 'tokenisation', 'crap', 'and', 'I', 'am',
'interested', 'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\nnis',
'handled', 'ok?', 'This', 'and', 'whether', 'it', 'properly', 'deals', 'with',
'websites', 'like', 'www.ucd.ie\\nand', 'email', 'addresses', 'like', 'mark.keane@ucd.email.ie.',
'', 'Also', 'other', 'weirdities', 'like\\nth', 'great', "O'Neill",
', 'and', 'like', 'M*A*S*H', 'which', 'should', 'be', 'maybe', '7', 'tokens?\n']
```

```
>>> stemmer = PorterStemmer()
```

```
>>> stemmer.stem(wd) for wd in splitrawtext
SyntaxError: invalid syntax
```

```
>>> [stemmer.stem(wd) for wd in splitrawtext]
['So,', 'thi', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put', 't
ogethr', 'to', 'check', 'on\\nth', 'tokenis', 'crap', 'and', 'I', 'am', 'interest',
'in', 'how', 'I.B.M.', 'or', 'the', 'U.S.A.', 'or', 'the', 'USA\\ni', 'handled',
'ok?', 'Thi', 'and', 'whether', 'it', 'properli', 'deal', 'with', 'websit', 'l
ike', 'www.ucd.ie\\nand', 'email', 'address', 'like', 'mark.keane@ucd.email.ie.',
'', 'Also', 'other', 'weirditi', 'like\\nth', 'great', "O'Neil", 'and', 'like', 'M
*A*S*H', 'which', 'should', 'be', 'mayb', '7', 'tokens?\n']
```

```
>>> stemmer = LancasterStemmer()
Traceback (most recent call last):
  File "<pyshell#10>", line 1, in <module>
    stemmer = LancasterStemmer()
NameError: name 'LancasterStemmer' is not defined
```

```
>>> stemmer = nltk.LancasterStemmer()
```

```
>>> [stemmer.stem(wd) for wd in splitrawtext]
['so,', 'thi', 'is', 'just', 'a', 'bunch', 'of', 'text', 'that', "i've", 'put',
'togethr', 'to', 'check', 'on\\nth', 'tok', 'crap', 'and', 'i', 'am', 'interest',
'in', 'how', 'i.b.m.', 'or', 'the', 'u.s.a.', 'or', 'the', 'usa\\nis', 'handled',
'ok?', 'thi', 'and', 'wheth', 'it', 'prop', 'deal', 'with', 'websit', 'lik', 'www
.ucd.ie\\nand', 'email', 'address', 'lik', 'mark.keane@ucd.email.ie.', '', 'also',
'oth', 'weird', 'like\\nth', 'gre', "o'neill", 'and', 'lik', 'm*a*s*h', 'which',
'should', 'be', 'mayb', '7', 'tokens?\n']
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
>>>
```

```
Ln: 23 Col: 0
```

```
MOBILEMAC:~ users
[Restored]
Last login: Sun Jun 15 12:53:38 on ttys000
MobileMac:~ users
[Restored]
Last login: Sun Jun 15 18:01:21 on ttys000
MobileMac:~ user$ idle
```

we did not import method

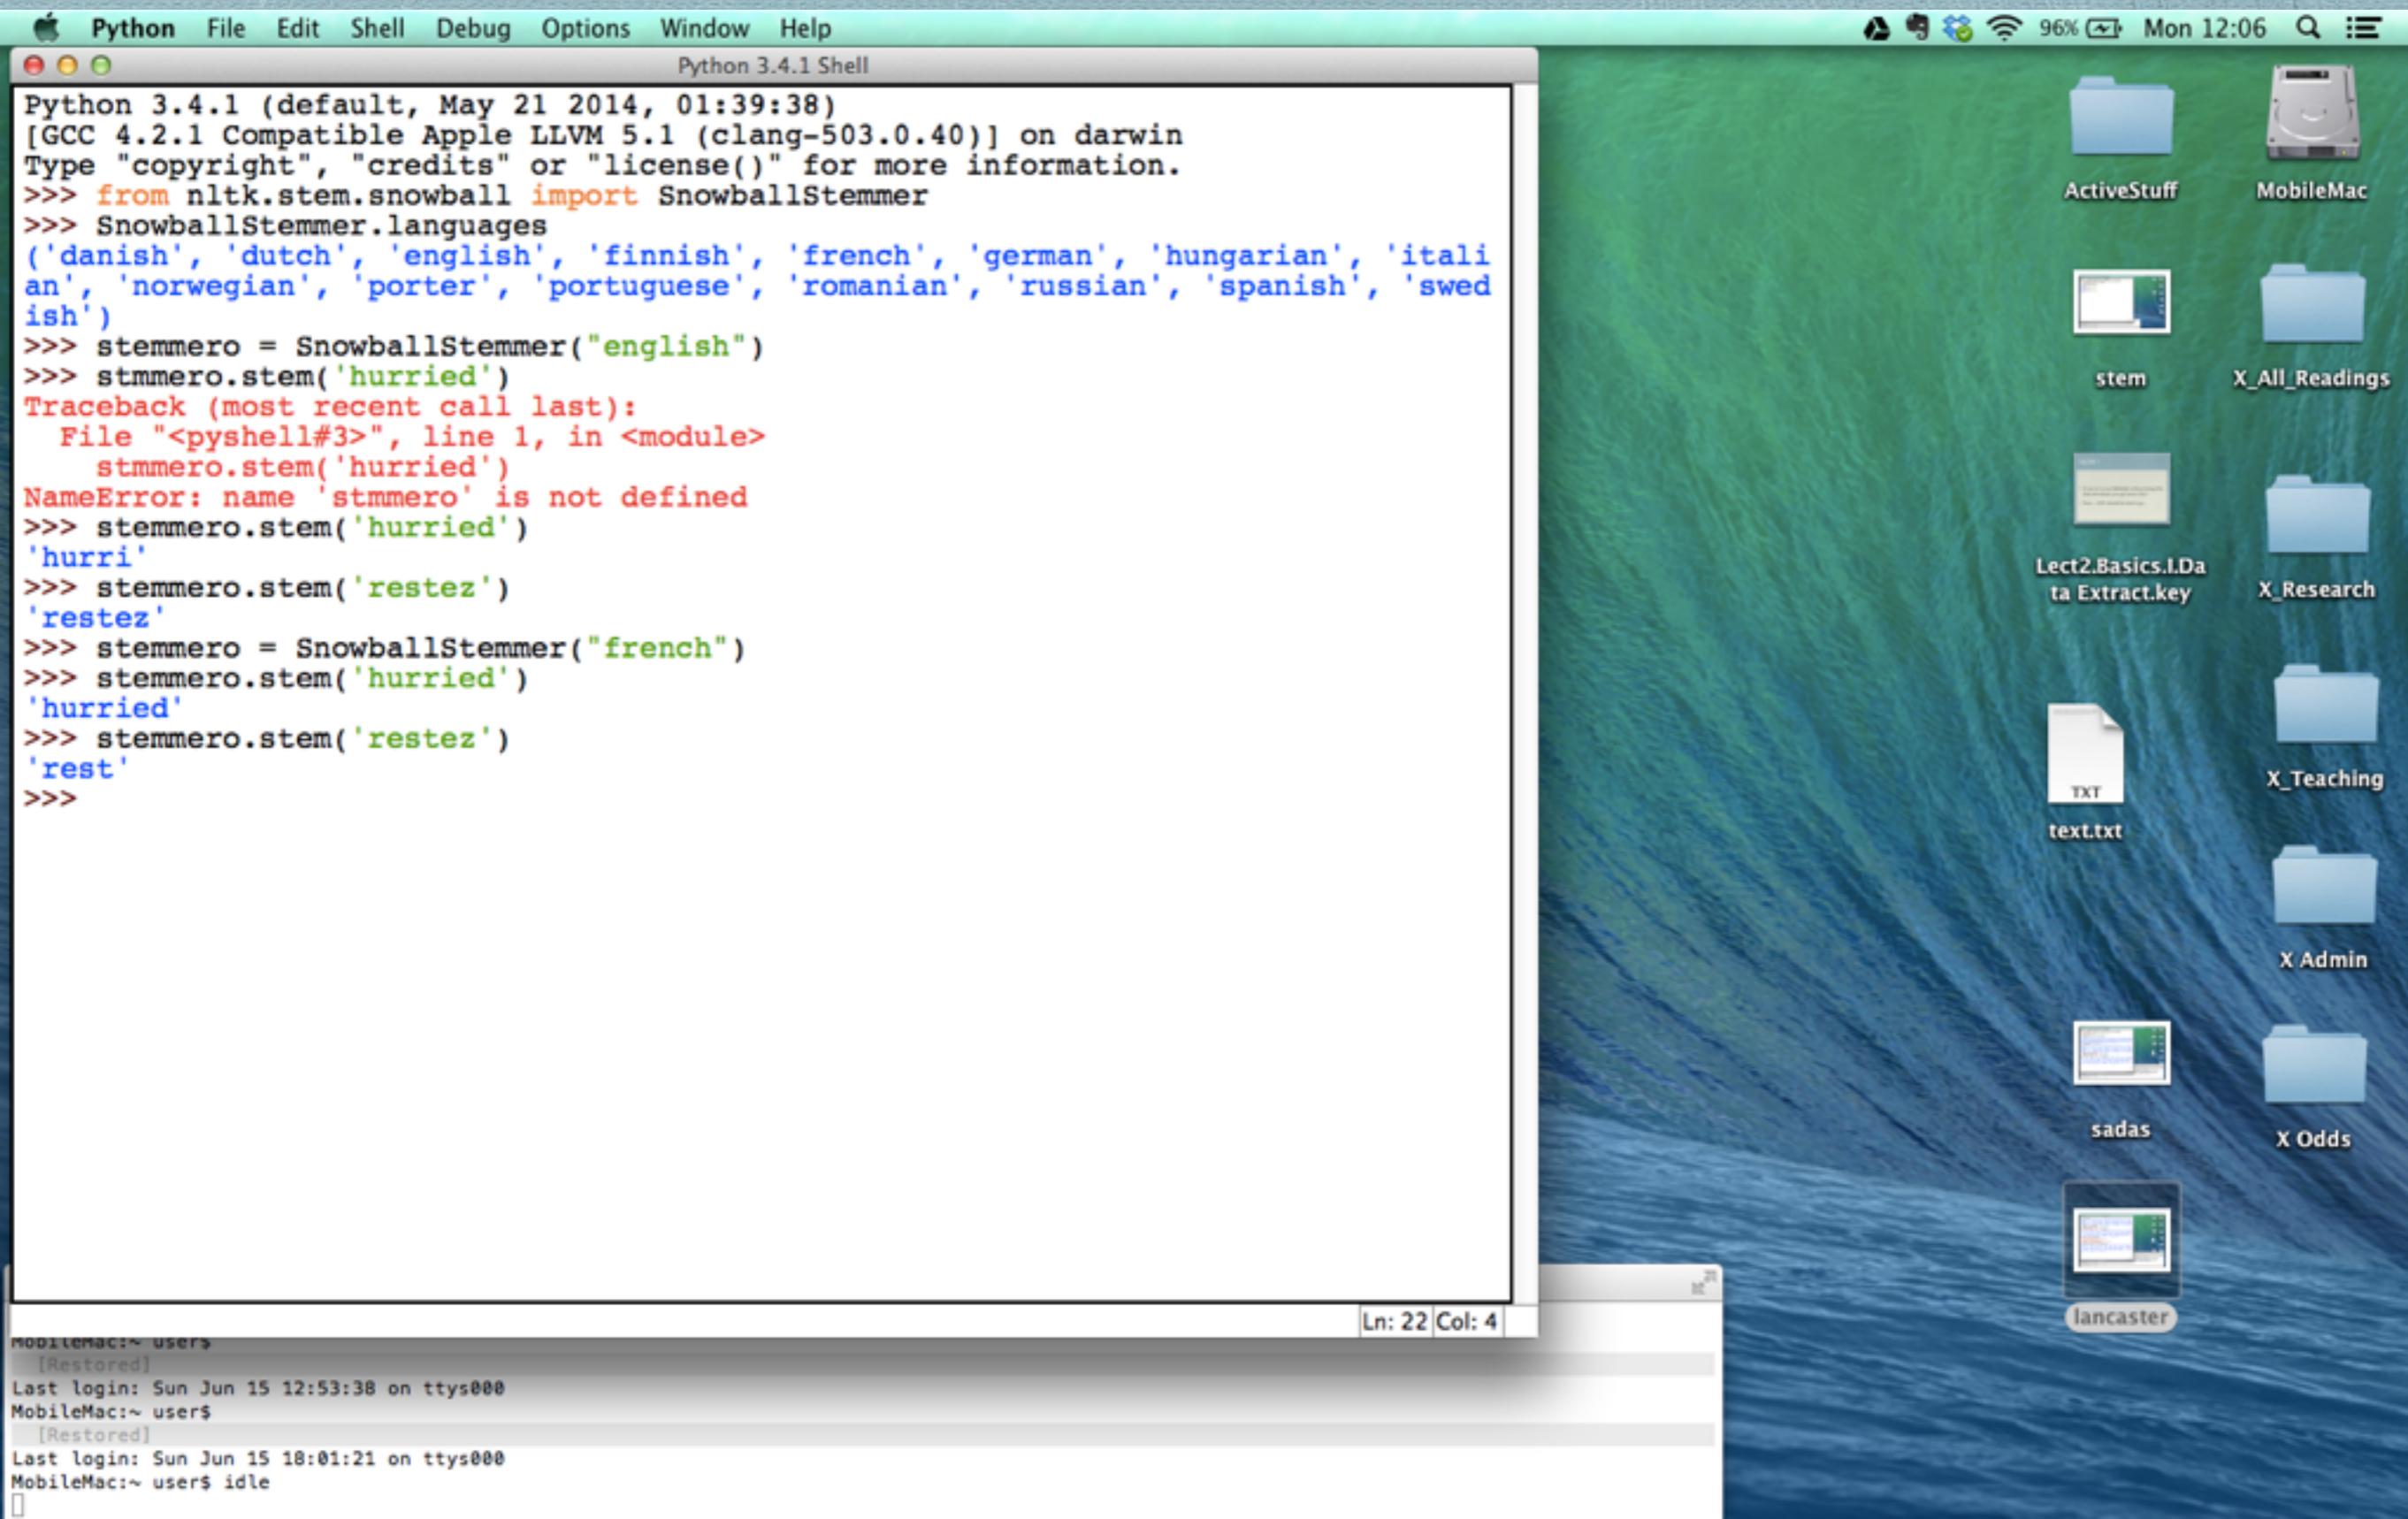
but, can name it explicitly instead of import

put togethr to check on
in how I.B.M. or the U.S.A. or the USA
ly deals with websites like www.ucd.ie
and email addresses like mark.keane@ucd.email.ie. Also other weirdities like
the great O'Neill and like M*A*S*H, which should be maybe 7 tokens?

Snowball Stemming

- ◆ `nltk` has a bunch of off-the-shelf stemmers
- ◆ Snowball Stemmer, in `nltk`, is really a language for building stemmers, as they are language-specific this is useful
- ◆ Also, you can use it to explicitly switch languages

Snowball Stemming: Eg



Butt Stemmers...But

- ◆ Are rather crude; no word meaning disambiguation (bats, batting)
- ◆ No POS disambiguation (fish and fish)
- ◆ Can't handle irregulars (am, is, are)
- ◆ So, sometimes you need more...

Why Lemmatise?

- ◆ Lemmatisation is RollsRoyce stemming
- ◆ Only removes affixes if the resulting word is in its dictionary (so, is slower...)
- ◆ So, it should be more accurate; see Wordnet (Miller, 1995), Sketch Engine (Kilgarriff et al, 2004)
- ◆ Some systems can get quite complex, do partial parses and contain frequency info for lemmas